

# Efficient Volume Estimation of Convex Polytopes in High Dimensions



Deepak, Raghu, Samuel, Vraj

Advisor: Yann Girsberger



Eidgenössische Technische Hochschule Zürich  
Swiss Federal Institute of Technology Zurich

# Algorithm




**Given Polytope description:**  $Q = \{x : Ax \leq b\}, A \in \mathbb{R}^{m \times n}, b \in \mathbb{R}^m$

**Task:** Find volume of above polytope

**Approach:**

1. Preprocess: Affine transform  $T$  s.t.  $B(0, 1) \subseteq P = T(Q) \subseteq B(0, r)$
2. EstimateVol: Concentric balls  $\{B_i\} = \{B(0, 2^{i/n})\}_{\forall i \in [\ell] \cup \{0\}}$

# Algorithm



$$\ell = \lceil n \log_2 r \rceil, K_i = B_i \cap P$$

$$\text{vol}(P) = \text{vol}(B(0, 1)) \prod_{i=0}^{\ell-1} (\text{vol}(K_{i+1}) / \text{vol}(K_i))$$

$$\alpha_i = \text{vol}(K_{i+1}) / \text{vol}(K_i)$$

Estimate  $\alpha_i$  using random walk (Monte Carlo Method)

# Cost Analysis

- Cost measure: adds, subs, mults, divs
- Preprocess: Ellipsoid method for rounding (polynomial time)
- Estimatevol: Calls `walk()`  $\leq 1600 (n \log_2 (2n))^2$  times
- Walk:  $(2mn + m + 2n + \mathcal{O}(1))$  flops per call.
- Overall complexity of Estimatevol is  $\tilde{\mathcal{O}}(mn^3)$
- As  $n$  increases, we expect the bottleneck to be Walk

# Baseline

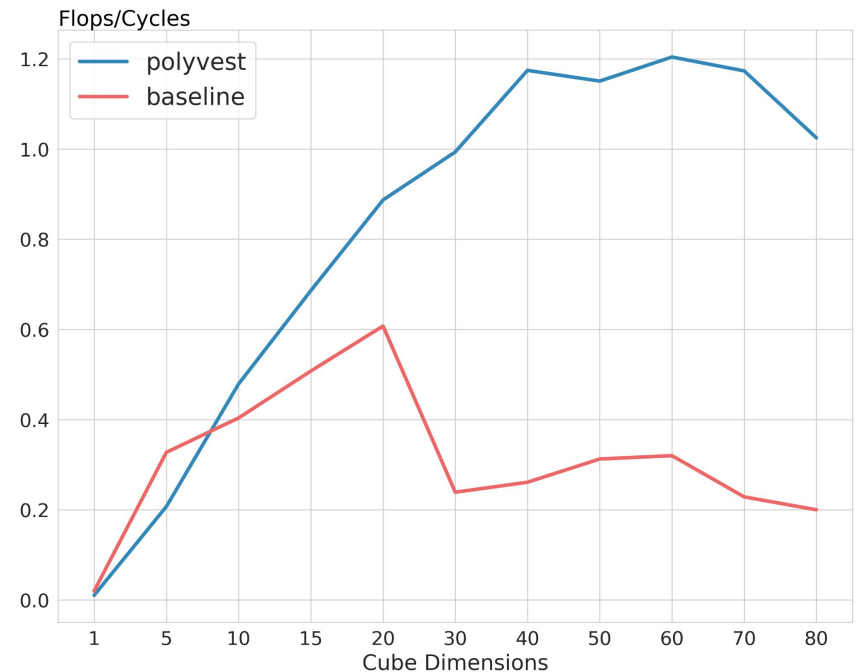
## Baseline:

- Focus on algorithm in [Ge et al. 2013](#)
- No additional optimization
- Benchmark against [PolyVest](#)

## Benchmark:

- Linux perf
- 5 repeats
- E.g.: `perf stat -M FLOPc -r 5 ./polyvol tests/cube_80`

All benchmarks were performed on:  
Intel Comet Lake i7-10610U @ 1.80GHz  
Compiler: GCC 11.3.0  
-O3 -march=native -ffast-math



Performance on cubes of varying dimensions

# Hotspot Analysis

- Basic Hotspot analysis (Intel VTune)
- Focus on polytope::walk

Function Stack ▼	CPU Time: Total »	CPU Time: Self »
Total	100.0%	0s
▼ _start	100.0%	0s
▼ __libc_start_main_impl	100.0%	0s
▼ main	100.0%	0s
▶ polytope::preprocess	0.0%	0s
▼ polytope::estimateVol	100.0%	7.817s
▶ polytope::walk	98.6%	369.062s
▶ func@0x2754	0.0%	0.020s
▶ func@0x2594	0.0%	0.008s
▶ __pow	0.3%	2.273s
▶ __log2	0.1%	0.978s
▶ __ceil_sse41	0.0%	0.072s
▶ [Unknown stack frame(s)]	0.0%	0s

**Hotspots for baseline on cube\_40**

# Optimization - I

- Used DP for more efficient computation of volume of a unit ball.
- Unroll loops in norm and estimatevol for ILP.
- Used more efficient random number generator (Xoshiro).
- Removed Bound checks from array access.
- Precomputed division operations in walk.

# Optimization - I

$A \in \mathbb{R}^{m \times n}$ ,  $B$  is an array of vectors.

---

**Algorithm 1:** Walk (Before pre-computation)

---

```
1  $d \leftarrow \text{RANDOM}(\{1, 2, 3, \dots, n\})$   
2 Constant Time computation  
3  $bound \leftarrow B[d] - (A / (A[:, d].\mathbf{1}^T)).x$  // Bottleneck -  $\mathcal{O}(m + mn)$  time  
4  $\mathcal{O}(m)$  computation for MIN and MAX  
5  $x_d \leftarrow x_d + \text{RANDOM}(\text{MIN}, \text{MAX})$  // Update x in the chosen direction
```

---



# Optimization - I

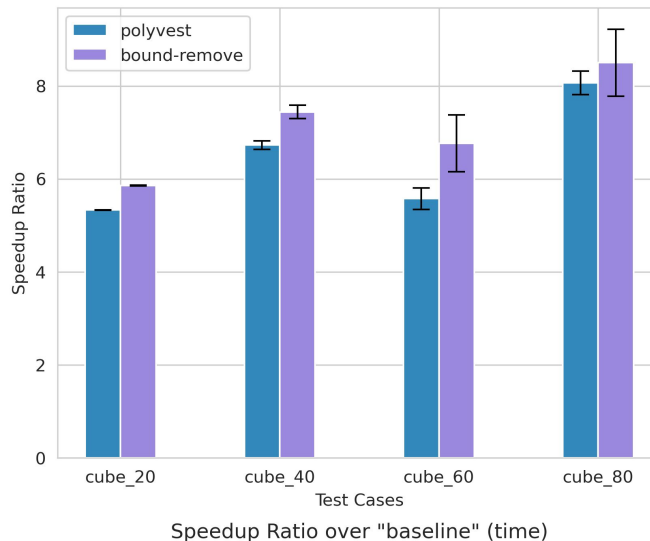
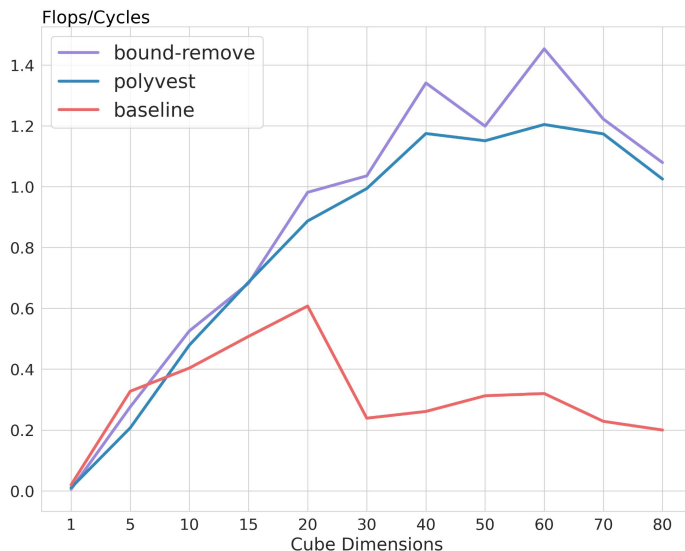
$$\forall j \in [n] \quad \tilde{A}[j] := A / (A[:, j] \cdot \mathbf{1}^T)$$

---

**Algorithm 2:** Fast Walk I: After pre-computation

---

- 1  $d \leftarrow \text{RANDOM}(\{1, 2, 3, \dots, n\})$
  - 2 Constant Time computation
  - 3  $\text{bound} \leftarrow B[d] - \tilde{A}[d].x$  // Bottleneck -  $\mathcal{O}(m + mn)$  time
  - 4  $\mathcal{O}(m)$  computation for MIN and MAX
  - 5  $x_d \leftarrow x_d + \text{RANDOM}(\text{MIN}, \text{MAX})$  // Update  $x$  in the chosen direction
- 



# Optimization - II

---

**Algorithm 3:** Fast Walk II: Asymptotically Fast

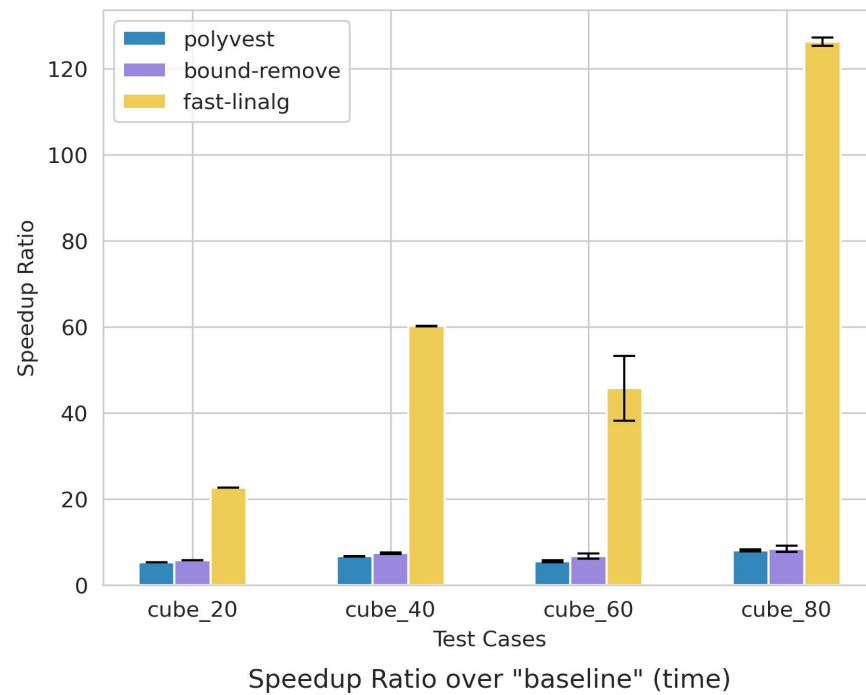

---

```
1  $d \leftarrow \text{RANDOM}(\{1, 2, 3, \dots, n\})$ 
2 Constant Time computation
3  $bound \leftarrow B[d] - A.x/A[:, d]$  // Now -  $\mathcal{O}(m)$  time
4  $\mathcal{O}(m)$  computation for MIN and MAX
5  $v \leftarrow \text{RANDOM}(\text{MIN}, \text{MAX})$ 
6  $x_d \leftarrow x_d + v$  // Update x in chosen direction
7  $A.x \leftarrow A.x + vA[:, d]$  // x changed, so update A.x -  $\mathcal{O}(m)$  time
```

---



Maintain the vector  $A.x$

# Optimization - II



# Optimization - III

- Vectorized norm computation and loops in `walk`
- Vectorized calculations in `estimateVol`.
- Precomputed some vector data types used in `walk`

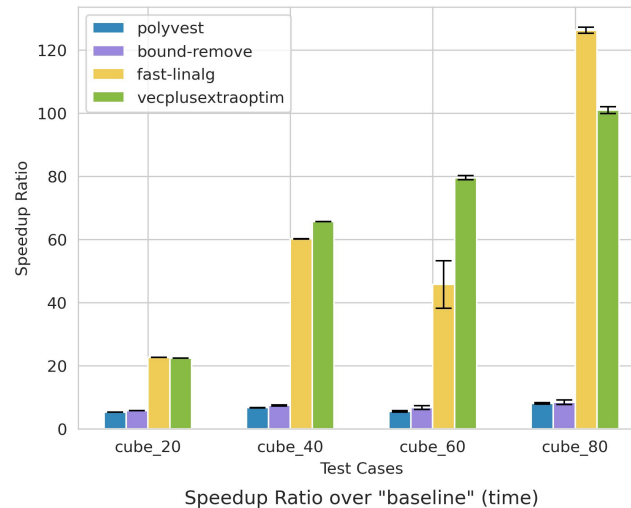
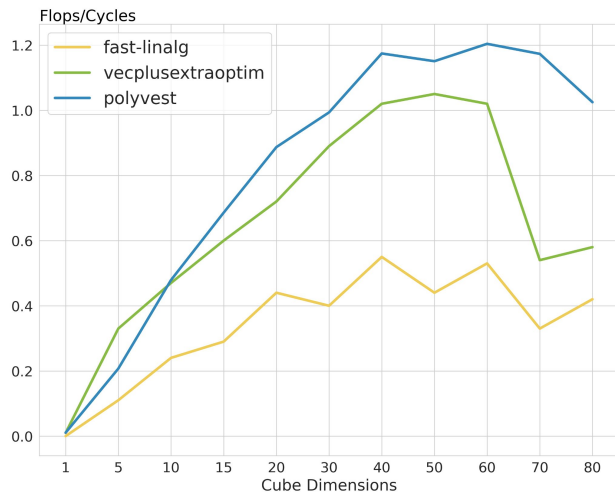


# Optimization - III

## Why no speedup for larger dimensions?

Function	Module	CPU Time ⓘ	% of CPU Time ⓘ
<code>_posix_memalign</code>	libc.so.6	98.187s	44.2%
<code>_ZNK8polytope4walkEPdSO_PKdRKN4arma3MatldEEPdV4_dSA_dRN10XoshiroCpp18Xoshiro128PlusPlusE</code>	polyvol	28.207s	12.7%
<code>__memcpy_avx_unaligned_erms</code>	libc.so.6	26.534s	11.9%
<code>__GI__</code>	libc.so.6	8.564s	3.9%
<code>_Z15_mm256_store_pdPdDv4_d</code>	polyvol	6.422s	2.9%
[Others]	N/A*	54.255s	24.4%

```
vec A_dir = A.col(dir), A_negrecp_dir = A_negrecp.col(dir);
```



# Optimization - IV

- Replaced most Armadillo objects with (aligned) double\*
  - Simpler Pointer Arithmetic
  - Further Vectorization possible
  - Column Major Format
- Reduced precision - From doubles to floats

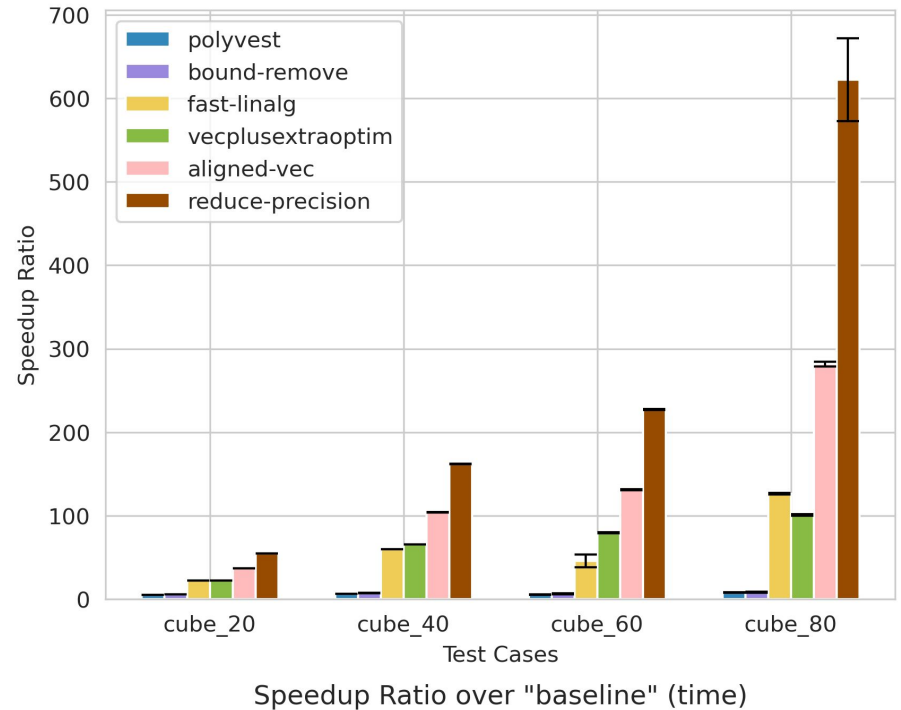

Function	Module	CPU Time ⓘ	% of CPU Time ⓘ
<a href="#">_ZNK8polytope4walkEPdS0_PKdS2_PKdV4_dS5_dRN10XoshiroCpp18Xoshiro128PlusPlusE</a>	polyvol	40.680s	34.8%
<a href="#">_mm256_blendv_pd</a>	polyvol	12.787s	10.9%
<a href="#">_Z15_mm256_store_pdPdDv4_d</a>	polyvol	7.890s	6.8%
<a href="#">_mm256_max_pd</a>	polyvol	7.676s	6.6%
<a href="#">norm_2</a>	polyvol	7.452s	6.4%
[Others]	N/A*	40.355s	34.5%

Allocations now outside Walk

```
- vec A_dir = A.col(dir), A_negrecp_dir = A_negrecp.col(dir);  
+ double* A_dir = A_arr + m*dir;
```

```
+ double *A_negrecp = (double *) aligned_alloc(32, align_pad(m*n));
```

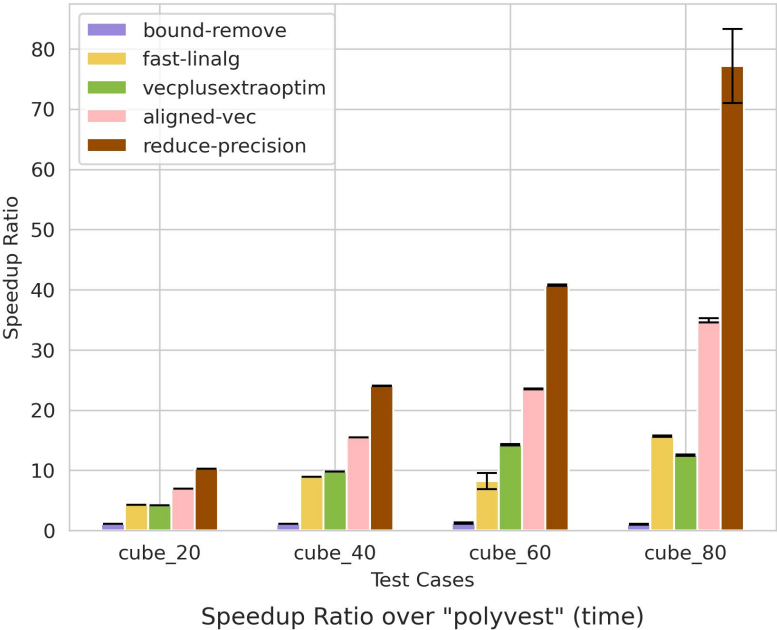
# Optimization - IV - Results



`aligned-vec` : Optimizations removing memory allocations, introducing aligned objects resulting in simpler pointer arithmetic and further vectorization

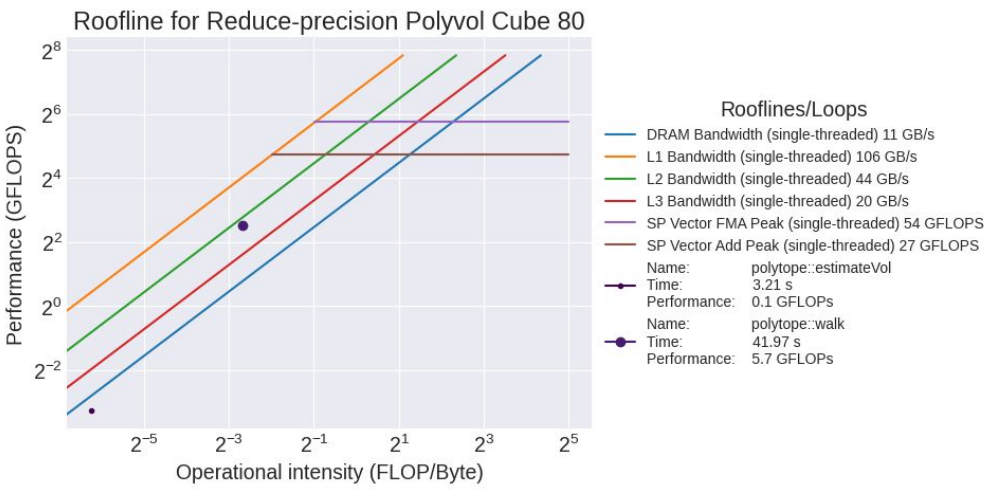
`reduce-precision` : Precision reduction from doubles to floats on top of `aligned-vec`

# Final Results



Method	Precision	GFLOPS	% Peak
Baseline	DP	0.175	0.62
Reduce-Precision	SP	4.742	8.28

Percentage Peak Performance





# Misc Optimizations

- **Profile-Guided Optimization for cube\_80**

Branch	PGO Time (s)	No PGO Time (s)
Vecplusextraoptim	228.65s $\pm$ 3.33	231.88s $\pm$ 1.29
Reduce-precision	45.23s $\pm$ 0.331	38.83s $\pm$ 5.46

- **Onefile**

No Improvement when including everything in one file.

- **Different Compiler (CLANG)**